

I will provide a picture of each function that I wrote. Each function has a description within the code comments.

Photos 1-5 are from the valueiterationagents.py file.

Photos 6-8 are from analysis.py

The rest are from qlearningagents.py

The very last picture is the autograder results.

Enjoy!

```
mdp.py    featureExtractors.py    valueIterationAgents.py

def __init__(self, mdp, discount = 0.9, iterations = 100):
    """
    Your value iteration agent should take an mdp on
    construction, run the indicated number of iterations
    and then act according to the resulting policy.

    Some useful mdp methods you will use:
    Note: these are described in mdp.py, defined in gridworld.py
    Note: these pull needed values from the mdp object
        mdp.getStates()
            Return a list of all states in the MDP.
        mdp.getPossibleActions(state)
            Returns list of actions from state
        mdp.getTransitionStatesAndProbs(state, action) t(s,a,s')
            Returns list of (nextState, prob) pairs
            Note: if terminal state returns []
        mdp.getReward(state, action, nextState)
            Returns reward from t(s,a,s')
        mdp.isTerminal(state)
            Returns true if the current state is a terminal state.
    """
    #compared to algorithm in textbook Fig 17.4
    self.mdp = mdp          # input mdp
    self.discount = discount # input gamma
    self.iterations = iterations # number of iterations defined by user
    self.values = util.Counter() # --U-- A Counter is a dict with default 0
    self.runValueIteration() # function VALUE-ITERATION call to run
```

```
mdp.py    featureExtractors.py    valueIterationAgents.py

def runValueIteration(self):
    # Write value iteration code here
    """ YOUR CODE HERE """
    """ This value iteration function checks that the state is not a
    terminal state and then runs a loop for each iteration of the algorithm
    calculating the maximum value of the possible actions and filling a vector
    with the q values. It calls getAction and getQValue which I define later.
    These two functions do all of the calculations necessary to determine the
    best action to take for an optimum policy. """

    # runs this loop (-i ...) specified in input
    for i in range(self.iterations): # for each iteration

        # get a new dict for this iteration U <- U'
        newStateValues = self.values.copy()

        s = self.mdp.getStates()
        # for each state in s do
        for state in s:

            terminal = self.mdp.isTerminal(state)
            # if the current state is NOT terminal then compute action/Qval
            if not terminal:
                action = self.getAction(state) # get max action
                newStateValues[state] = self.getQValue(state, action)

        # re populate the original dict with the new values
        self.values = newStateValues

    # completed value iteration
```

```
#these two functions just call the functions I defined.
def getAction(self, state):
    # "Returns the policy at the state (no exploration)."
    return self.computeActionFromValues(state)

def getQValue(self, state, action):
    return self.computeQValueFromValues(state, action)

def computeQValueFromValues(self, state, action):
    """
    Compute the Q-value of action in state from the
    value function stored in self.values.
    """
    *** YOUR CODE HERE ***
    """
    This function uses the Bellman equation to calculate the utility
    of a q-state. First we get the state (called nextState) that is
    reachable by taking an action from the given state along with
    it's probability of occurring. Next we find the reward of going to
    that next state. Next, we calculate the utility using the Bellman
    equation to find the utility (currentStateValue) of k+1 for the
    input state.
    """

    currentStateValue = 0 # value starts at zero

    for nextState, p in self.mdp.getTransitionStatesAndProbs(state, action):
        r = self.mdp.getReward(state, action, nextState) # get reward
        y = self.discount # get discount (gamma)
        nextU = self.getValue(nextState) # get the max utility of next state
        currentStateValue += p * (r + y * nextU) # calculate U of k+1 given s
    return currentStateValue
```

```
def computeActionFromValues(self, state):
    """
    The policy is the best action in the given state
    according to the values currently stored in self.values.

    You may break ties any way you see fit. Note that if
    there are no legal actions, which is the case at the
    terminal state, you should return None.
    """
    """ YOUR CODE HERE """
    """
    This function checks if the state is terminal and if it is
    then returns None as in the instructions. Next, it calculates the
    values of each action and stores them in a temporary dict for
    comparison. Next, the function argMax() defined in util.py is used
    to find the maximum value to be returned which is eventually
    multiplied with gamma in the Bellman equation.
    """
    # if there aren't any legal actions, return None (break ties)
    if self.mdp.isTerminal(state):
        return None

    # create a temp dict for comparing actions
    tempCache = util.Counter()

    # get the optimum policy value of each state/action/Qval
    for action in self.mdp.getPossibleActions(state):
        tempCache[action] = self.getQValue(state, action)
    # use argmax to find best policy
    # (returns key with highest value defined in util.py line 334)
    policy = tempCache.argmax()
    return policy # return the best action from the given state
```

```
mdp.py    featureExtractors.py    analysis.py    valueIterationAgents.py

#####
# ANALYSIS QUESTIONS #
#####

# Set the given parameters to obtain the specified policies through
# value iteration.

"""      Question Two

It seems that any value less than .016 for the noise parameter causes the
probability of taking the action east starting from state (1,1) and endings
in state (6,1) to become positive. The original value (0.02) for noise causes
a negative probability in state (1,1) of going east. As I tested noise values,
I decreased the noise parameter by 0.005 until I started to see changes.
Slowly I watched the far right states turn green, but they all had to become
green for a reward of ten.
"""

def question2():
    answerDiscount = 0.9
    answerNoise = 0.016 # 0.2
    return answerDiscount, answerNoise
```

```
""" Question 3
I found these mainly by trial and error. You can watch the affects of each try
on the different q values for taking an action from a state and how it changes.
"""

def question3a():
    # You can use a very small value of gamma to cause the bellman equation to
    # add the value of the next states reward. Since we get to the rewards
    # faster by going right, this method works for our wanted policy type.
    answerDiscount = .1 # or even .000000001
    answerNoise = 0
    answerLivingReward = 0
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'

def question3b():
    answerDiscount = .2
    answerNoise = .2
    answerLivingReward = .5
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'

def question3c():
    answerDiscount = .9
    answerNoise = .02
    answerLivingReward = 0
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'

def question3d():
    answerDiscount = .9
    answerNoise = .5
    answerLivingReward = 0
    return answerDiscount, answerNoise, answerLivingReward
```

```
def question3d():
    answerDiscount = .9
    answerNoise = .5
    answerLivingReward = 0
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'

def question3e():
    answerDiscount = .9
    answerNoise = .02
    answerLivingReward = 1
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'

def question8(): #after testing values, I determined that it is not possible.
    answerEpsilon = None
    answerLearningRate = None
    return 'NOT POSSIBLE'
```

```
def __init__(self, **args):
    "You can initialize Q-values here..."
    ReinforcementAgent.__init__(self, **args)

    """ YOUR CODE HERE """
    # Initialize the q values by assigning a new dict to it
    # A Counter is a dict with default 0
    self.Q_value = util.Counter()

def getQValue(self, state, action):
    """
    Returns Q(state,action)
    Should return 0.0 if we have never seen a state
    or the Q node value otherwise
    """
    """ YOUR CODE HERE """
    # This function returns 0 if there is no possible actions, this means we
    # haven't seen a state yet. If there is possible actions we return the
    # Q node value.

    if not self.getLegalActions(state):
        return 0.0
    else:
        return self.Q_value[(state, action)]
```



```
def computeValueFromQValues(self, state):  
    """  
    Returns max_action Q(state,action)  
    where the max is over legal actions. Note that if  
    there are no legal actions, which is the case at the  
    terminal state, you should return a value of 0.0.  
    """  
    """ YOUR CODE HERE """  
  
    legalActions = self.getLegalActions(state)  
  
    #if no legal actions return 0.0  
    if not legalActions:  
        return 0.0  
  
    #otherwise we assign value to a really small negative number as a default  
    else:  
        maxValue = float('-inf')  
  
    #next we update value to the highest value from all states within reach  
    for action in legalActions: # for each action  
        Qval = self.getQValue(state, action) #get Q value of that action  
        maxValue = max(maxValue, Qval) #take max of previous value and Qval  
    return maxValue #finally return the max value
```

```
def computeActionFromQValues(self, state):  
    """  
        Compute the best action to take in a state. Note that if there  
        are no legal actions, which is the case at the terminal state,  
        you should return None.  
    """  
    """ YOUR CODE HERE """  
    legalActions = self.getLegalActions(state)  
  
    #if no legal actions return None  
    if not legalActions:  
        return None  
  
    bestAction = [] # no best action yet make blank  
  
    #get the action with the max value from the rest of the Qnodes  
    QValue = self.computeValueFromQValues(state)  
  
    # for each legal action  
    for action in legalActions:  
        Qval = self.getQValue(state, action) #get Qvalue of this state/action  
        if QValue == Qval: #if they match add to list of best actions  
            bestAction.append(action)  
    return random.choice(bestAction) # choose a best action at random
```

```
def getAction(self, state):  
    """  
    Compute the action to take in the current state. With  
    probability self.epsilon, we should take a random action and  
    take the best policy action otherwise. Note that if there are  
    no legal actions, which is the case at the terminal state, you  
    should choose None as the action.  
    HINT: You might want to use util.flipCoin(prob)  
    HINT: To pick randomly from a list, use random.choice(list)  
    """  
    epsilon = self.epsilon #define epsilon for this function  
    legalActions = self.getLegalActions(state) #list of possible actions  
  
    #if no legal actions return None  
    if not legalActions:  
        return None  
    else:  
        # with a small probability we act randomly  
        if util.flipCoin(epsilon): # with a probability of epsilon  
            return random.choice(legalActions) # we return a random action  
  
        # or we return the best action with probability of 1 - epsilon  
        else:  
            return self.computeActionFromQValues(state)
```

```
def update(self, state, action, nextState, reward):
    """
    The parent class calls this to observe a
    state = action => nextState and reward transition.
    You should do your Q-Value update here
    NOTE: You should never call this function,
    it will be called on your behalf
    """
    """ YOUR CODE HERE """
    """ Update's the Q-Value (Qk+1) I designed this to match slide 57
    """

    gamma = self.discount    #discount
    alpha = self.alpha        #learning rate
    max_qk_sPrime_a = self.computeValueFromQValues(nextState)
    sample = reward + (gamma * max_qk_sPrime_a) #new sample
    qk_sa = self.Q_value[(state, action)]    #old sample (new sample later)

    # Incorporate the new estimate into a running average
    self.Q_value[(state, action)] = qk_sa + alpha * (sample - qk_sa)

    # or this is equivalent
    # self.Q_value[(state, action)] = ((1 - alpha) * qk_sa) + (alpha * sample)

def getQValue(self, state, action):
    """
    Should return Q(state,action) = w * featureVector
    where * is the dotProduct operator
    """
    """ YOUR CODE HERE """
    #  $Q(s,a) = w_1f_1(s,a) + w_2f_2(s,a) + \dots + w_nf_n(s,a)$ 
    # this function calculator the  $w_nf_n(s,a)$  given the state/action pair
    # get weights and the vector of features and return product

    w = self.getWeights() # get the weight
    f = self.featExtractor.getFeatures(state, action) # get the feature
    q_s_a = w * f # compute the product
    return q_s_a # return the product
```

```
def update(self, state, action, nextState, reward):
    """
        Should update your weights based on transition
    """
    """*** YOUR CODE HERE ***"""
    # This function updates the weights of active features for the qlearner!
    # I wrote this to match the linear Q-functions on slide 68
    # The comments describe what part of the algorithm the code represents
    gamma = self.discount    #discount
    alpha = self.alpha        #learning rate
    # max_qk_sPrime_aPrime == maxQ(s',a')
    max_qk_sPrime_aPrime = self.computeValueFromQValues(nextState)
    stuff = reward + (gamma * max_qk_sPrime_aPrime) # == [r + y(maxQ(s',a'))]
    q_sa = self.getQValue(state, action)    # == Q(s,a)
    difference = stuff - q_sa # == [r + y(maxQ(s',a'))] - Q(s,a)
    features = self.featurer.getFeatures(state, action) #get the features

    # for each feature update the weights
    for f in features:
        self.weights[f] += alpha * difference * features[f]
    return
```

Finished at 22:13:25

Provisional grades

=====

Question q1: 4/4

Question q2: 1/1

Question q3: 5/5

Question q4: 0/1

Question q5: 0/3

Question q6: 4/4

Question q7: 2/2

Question q8: 1/1

Question q9: 1/1

Question q10: 3/3

Total: 21/25

Note: The only reason why q4 and q5 are at a zero value is because they were not presented to me within the lab questions, so I had no way to answer them.